The AI "GPU Tax" Problem

The hidden infrastructure cost that prevents global AI from scaling efficiently.

The Invisible Barrier to Global AI Scaling

The Problem: Current LLM incumbents pay a full "GPU Tax" for every query, treating a 50-token greeting the same as a 2000-token research task.

The Impact: This makes high-performance AI economically unsustainable for the "Data-Poor" connected users in emerging markets.

50
Greeting tokens

2000
Research tokens

1x GPU
Same cloud billing path

Standard AI Cloud Costs

50-token greeting

500-token task

2000-token research

AI for All Orchestration Model

Edge / local first

Regional optimize next

GPU only when needed

Token-Aware Orchestration

Core architecture: route task complexity to the right compute tier instead of defaulting every query to cloud GPUs.

On-Device Tier

Token band: ≤400 tokens

Query mix: ~70% of common queries handled locally

Economics: Cloud cost: $0

Regional Optimized Tier

Token band: 400–1000 tokens

Query mix: ~20–25% of mid-range tasks

Economics: Low-latency, low-cost regional infrastructure

LLM Fallback Tier

Token band: ≥1000 tokens

Query mix: Only 5–10% of queries escalate

Economics: High-cost cloud GPUs used only when necessary

The Unit Economics of "AI for All"

A structurally different cost model for global AI access.

90%

Queries routed away from expensive cloud GPUs

5–10%

Only the hardest tasks hit cloud LLM fallback

Patent-Pending

Core edge-cloud orchestration logic (India Provisional)

Achieving a 90% Efficiency Advantage

The Moat: By routing 90% of queries away from expensive cloud GPUs, we fundamentally change the unit economics of AI.

Status: Core Edge-Cloud Orchestration logic is Patent-Pending (India Provisional).

Vision: Building the infrastructure for the next billion users, not just the top 1%.